

Best Practices with Terraform Scripts

Robert Rozas Navarro
Premier Field Engineer
Apps Domain



Agenda

1. Best Practices
2. Hands on Lab 2
3. Q&A




Terraform

Best Practices


Best Practice 1

- Use remote backend
- Your laptop is no place for your infrastructure source of truth
- Use `data sources` and `terraform_remote_state` specifically as a glue between infrastructure modules within composition
- Managing a `tfstate` file in git is a nightmare
- Later when infrastructure layers starts to grow in any direction (number of dependencies or resources)



Best Practice 1 (..Continued..)

- Using the backend functionality has definitely benefits:
- Working in a team: it allows for collaboration, the remote state will always be available for the whole team
- The state file is not stored locally. Possible sensitive information is now only stored in the remote state
- Some backends will enable remote operations. The terraform apply will then run completely remote. These are called the enhanced backends



Best Practice 1 (..Continued..)

- You can also store your state in Azure Storage:

```
terraform {  
  backend "azurerm" {  
    resource_group_name = "StorageAccount-ResourceGroup"  
    storage_account_name = "abcd1234"  
    container_name       = "tfstate"  
    key                   = "prod.terraform.tfstate"  
  }  
}
```

Best Practice 1 (..Continued..)

- When using an Azure remote state, it's best to configure the Service Principal credentials:

```
$ az ad sp create-for-rbac --role="Contributor" --scopes="/subscriptions/SUBSCRIPTION_ID"
```

This command will output 5 values:

```
{
  "appId": "00000000-0000-0000-0000-000000000000",
  "displayName": "azure-cli-2017-06-05-10-41-15",
  "name": "http://azure-cli-2017-06-05-10-41-15",
  "password": "0000-0000-0000-0000-000000000000",
  "tenant": "00000000-0000-0000-0000-000000000000"
}
```


Best Practice 2

- Manage Terraform, Azure provider and modules version
- While individual resources are like atoms in the infrastructure, resource modules are molecules. Module is a smallest versioned and shareable unit.
- Examples and Terraform modules should contain documentation explaining features and how to use them.
- Infrastructure modules and compositions should persist their state in a remote location which can be reached by others in a controllable way



Best Practice 2 (..Continued..)

- You can use modules to make your terraform project more organized
- Infrastructure modules and compositions should persist their state in a remote location which can be reached by others in a controllable way
- Use third party modules (i.e from github)
- Allows you to reuse parts of your code



Best Practice 2 (..Continued..)

Use a module from git

```
module "module-example" {  
    source = "github.com/wardviaene/terraform-module-example"  
}
```

Use a module from a local folder

```
module "module-example" {  
    source = "./module-example"  
}
```

Best Practice 2 (..Continued..)

Pass arguments to the module

```
module "module-example" {  
  source = "./module-example"  
  region = "us-west-1"  
  ip-range = "10.0.0.0/8"  
  cluster-size = "3"  
}
```

Best Practice 2 (..Continued..)

Inside the module folder, you just have again, terraform files:

module-example/vars.tf

```
variable "vm_os_simple" {
    default = ""
}


# Definition of the standard OS with "SimpleName" = "publisher,offer,sku"
variable "standard_os" {
    default = {
        "UbuntuServer" = "Canonical,UbuntuServer,16.04-LTS"
        "WindowsServer" = "MicrosoftWindowsServer,WindowsServer,2016-Datacenter"
        "RHEL"          = "RedHat,RHEL,7.5"
        "openSUSE-Leap" = "SUSE,openSUSE-Leap,42.2"
        "CentOS"        = "OpenLogic,CentOS,7.6"
        "Debian"         = "credativ,Debian,8"
        "CoreOS"         = "CoreOS,CoreOS,Stable"
        "SLES"           = "SUSE,SLES,12-SP2"
    }
}
```

module-example/cluster.tf

```
resource "azurerm_virtual_machine" "vm-linux" {
    count                = ! contains(list(var.vm_os_simple, var.vm_os_offer), "Windows")
    name                 = "${var.vm_hostname}${count.index}"
    location             = var.location
    resource_group_name = azurerm_resource_group.vm.name
    availability_set_id  = azurerm_availability_set.vm.id
    vm_size              = var.vm_size
    network_interface_ids = [element(azurerm_network_interface.vm.*.id, count.index)]
    delete_os_disk_on_termination = var.delete_os_disk_on_termination
}
```


Best Practice 3

- Use implicit dependencies
- Implicit dependencies should be used whenever possible (see [this article](#) from terraform.io website for more information).
- With IaC the resources will be configured exactly as declared, and implicit dependencies can be used to ensure the creation order.



Best Practice 4


- Try to practice a consistent structure and **naming convention**
- Use _ (underscore) instead of - (dash) in all: resource names, data source names, variable names, outputs.
- Only use lowercase letters and numbers.



Generating Images

- The output of terraform graph is in the DOT format, which can easily be converted to an image by making use of dot provided by GraphViz:

```
$ terraform graph | dot -Tsvg > graph.svg
```



How should I structure my Terraform configurations?

- What is the complexity of your project?
- How often does your infrastructure change?
- How environments are grouped?
- Do not store all your code in a single file
- Provide a file structure that allows “separation of concerns”
- You can try more than one way to do it.



Terraform configurations (..Continued..)

provider.tf

```
provider "azurerm" {  
  # Whilst version is optional, we /strongly recommend/ using it  
  version = "=1.38.0"  
  
  subscription_id = "00000000-0000-0000-0000-000000000000"  
  client_id       = "00000000-0000-0000-0000-000000000000"  
  client_secret   = "${var.client_secret}"  
  tenant_id      = "00000000-0000-0000-0000-000000000000"  
}
```

terraform.tfvars

```
$ export ARM_CLIENT_ID="00000000-0000-0000-0000-000000000000"  
$ export ARM_CLIENT_SECRET="00000000-0000-0000-0000-000000000000"  
$ export ARM_SUBSCRIPTION_ID="00000000-0000-0000-0000-000000000000"  
$ export ARM_TENANT_ID="00000000-0000-0000-0000-000000000000"
```

vars.tf

```
variable "resource_group_name" {  
  description = "The name of the resource group in which the resources v  
  default     = "terraform-compute"  
}  
  
variable "location" {  
  description = "The location/region where the virtual network is create  
}  
  
variable "vnet_subnet_id" {  
  description = "The subnet id of the virtual network where the virtual  
}
```

instance.tf

```
resource "azurerm_virtual_machine" "vm-linux" {  
  count                                = ! contains(list(var.vm_os_simple, var.vm_os_offer), "Windows")  
  name                                = "${var.vm_hostname}${count.index}"  
  location                             = var.location  
  resource_group_name                  = azurerm_resource_group.vm.name  
  availability_set_id                  = azurerm_availability_set.vm.id  
  vm_size                              = var.vm_size  
  network_interface_ids                = [element(azurerm_network_interface.vm.*.id, count.index)]  
  delete_os_disk_on_termination       = var.delete_os_disk_on_termination
```

Terraform configurations (..Continued..)

- Putting all code in main.tf is a good idea when you are getting started or writing an example code
- In all other cases you will be better having several files split logically like this:
 - main.tf - call modules, locals and data-sources to create all resources
 - variables.tf - contains declarations of variables used in main.tf
 - outputs.tf - contains outputs from the resources created in main.tf
- terraform.tfvars should not be used anywhere except composition.

Hands On Lab 1

Q&A?